



GNOMON<sup>®</sup>  
DIGITAL


*LLM*

Jiqiong QIU

# Objectif du cours



- Le but de ce cours est l'introduction aux modèles LLM (large language model)
- Le cours se focalise essentiellement sur les aspects suivants :
  - *Introduction à LLM*
  - *Prompt engineering*
  - *Semantic search*
  - *Longchain/ LlamaIndex*
  - *RAG*
  - *Function calling*





- Plus besoin de connaissance ML
- Pas besoin d'apprentissage
- Pas besoin des données
- Just besoin de bien formuler ton demande (prompt)





# LLM



GNOMON<sup>®</sup>  
DIGITAL

## AI Builders

LE LEADER DE LA STRATÉGIE DATA ET IA

### LLM Providers Landscape

#### LLM Commerciaux



#### LLM Open Source



#### LLM Privés



Generative language models


Input: Text  
Output: Text

Generative image models


Input: Text  
Output: Image

Multimodal:

Input: Text, audio, images, vidéo  
Output: Text, audio, images, vidéo



# LLM








Comment ça marche?

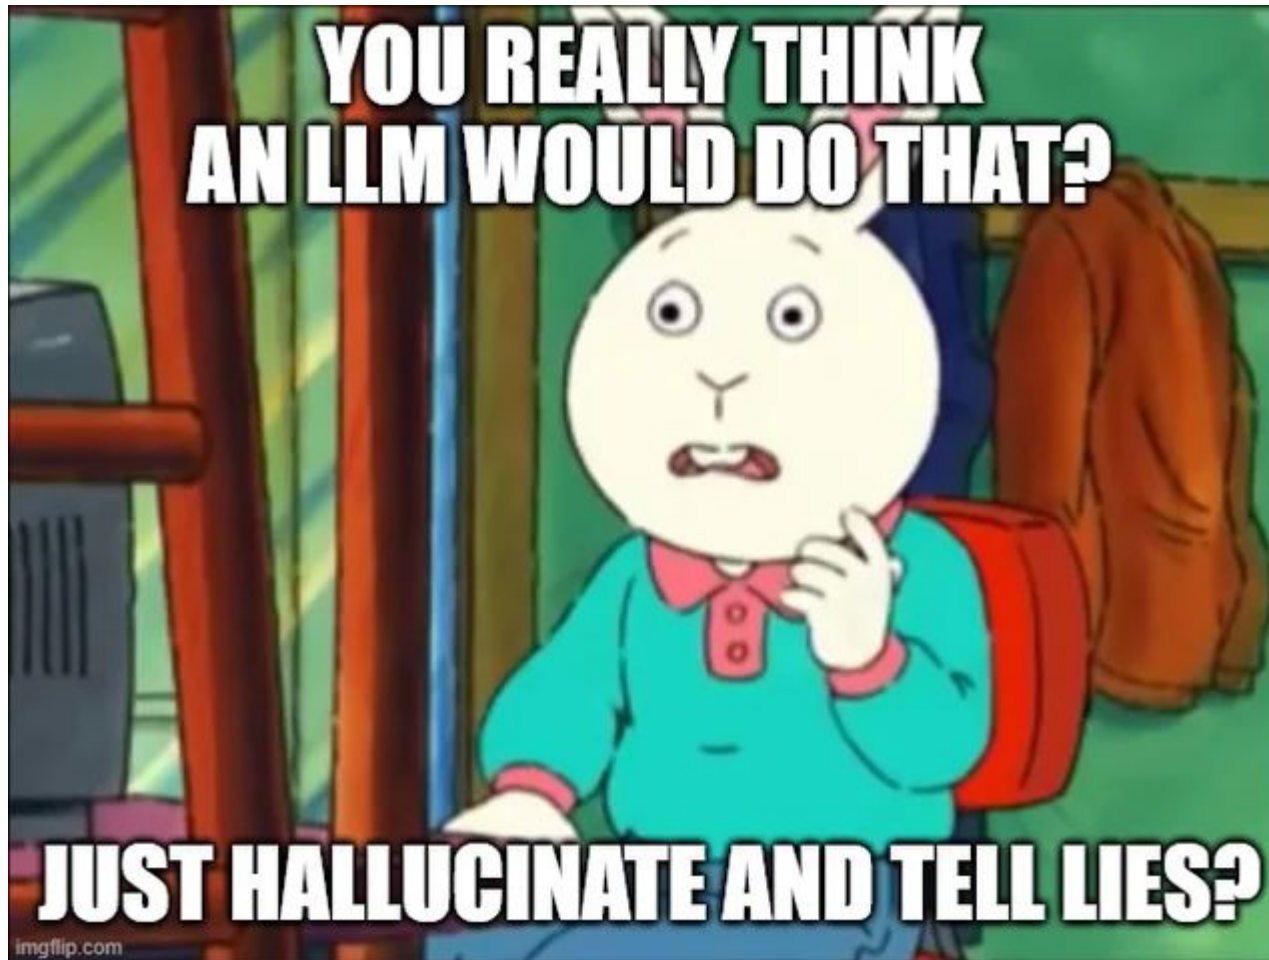
- Grand volume des données
- Billions parameters
- Apprentissage non-supervisé



# LLM



GNOMON<sup>®</sup>  
DIGITAL



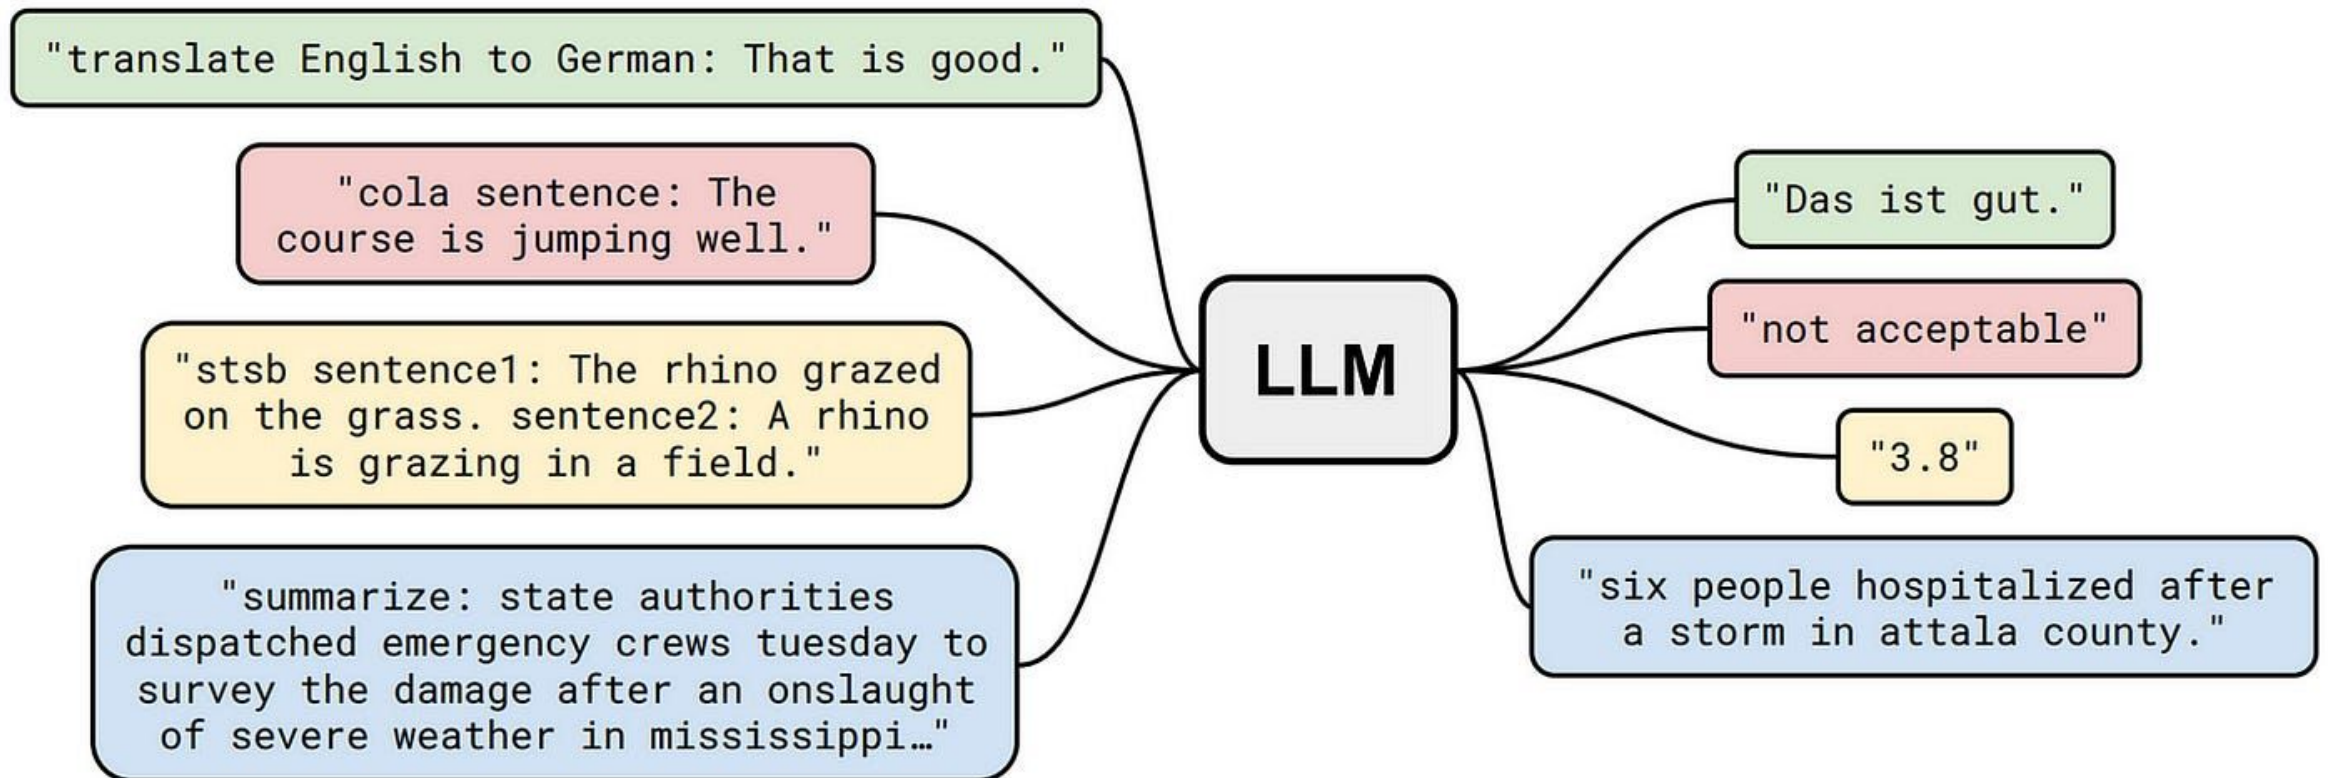
- Hallucinations:
  - Le modèle n'a pas assez des données pour faire l'entraînement
  - Le modèle est entraîné sur des données sales avec des bruits
  - On manque des context
  - Pas assez de contraintes sur le modèle

# Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

- On donne la context au modèle



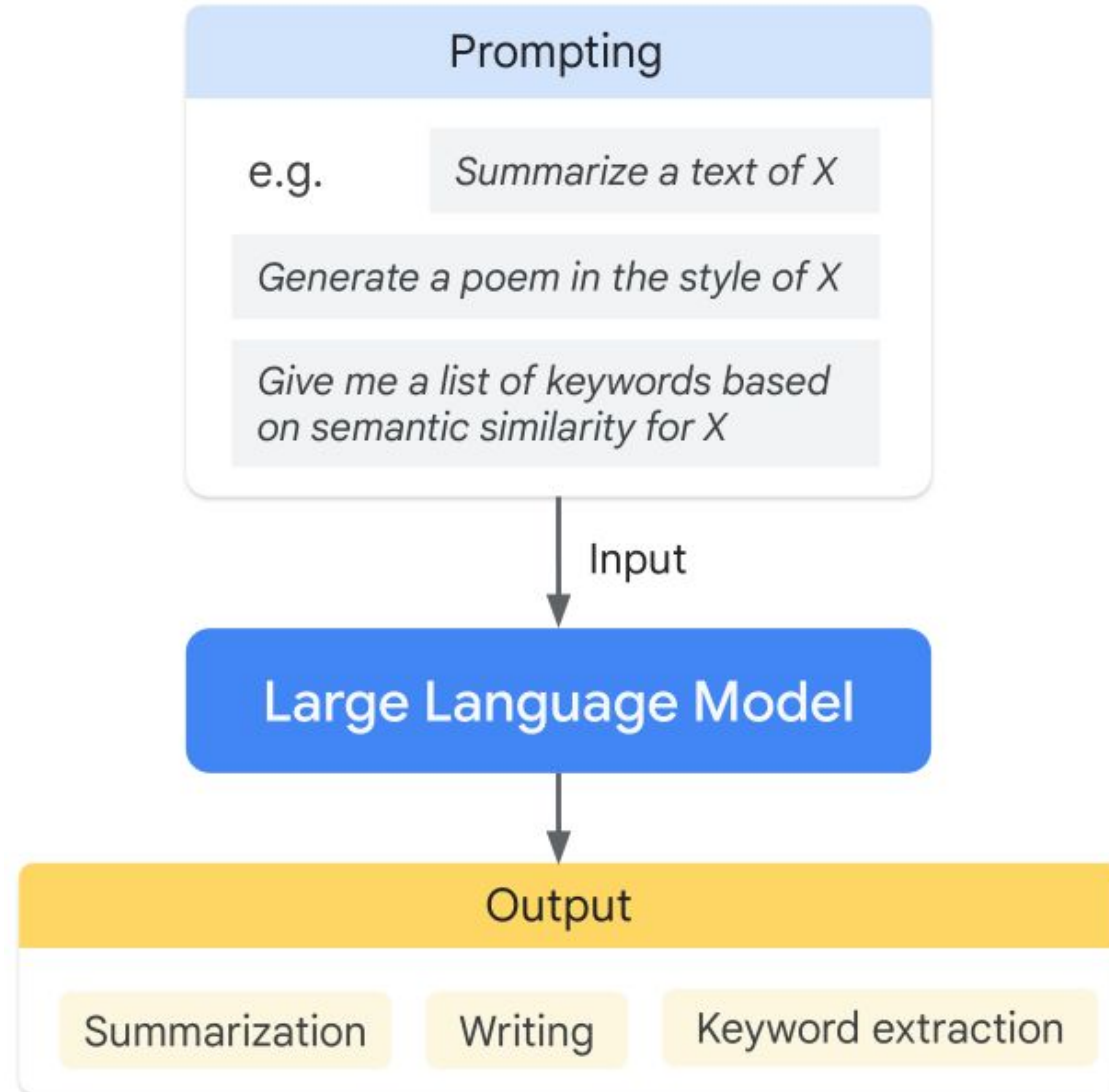


# Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

- Prompt engineering: la qualité d'entrée détermine la qualité de sortie



- Comment détermine un LLM



Forming a Database



Inputting a Prompt




Generating content

# LLM




GNOMON<sup>®</sup>  
DIGITAL



# LLM: type de modèles

- Comment détermine un LLM

Text-to-text



# LLM: type de modèles




GNOMON<sup>®</sup>  
DIGITAL

- Comment détermine un LLM

## Text-to-text

Domain d'application:

- Génération des text
- Classification
- Résumé
- Traduction
- Reformuler
- Extraction
- Clustering





# LLM: type de modèles

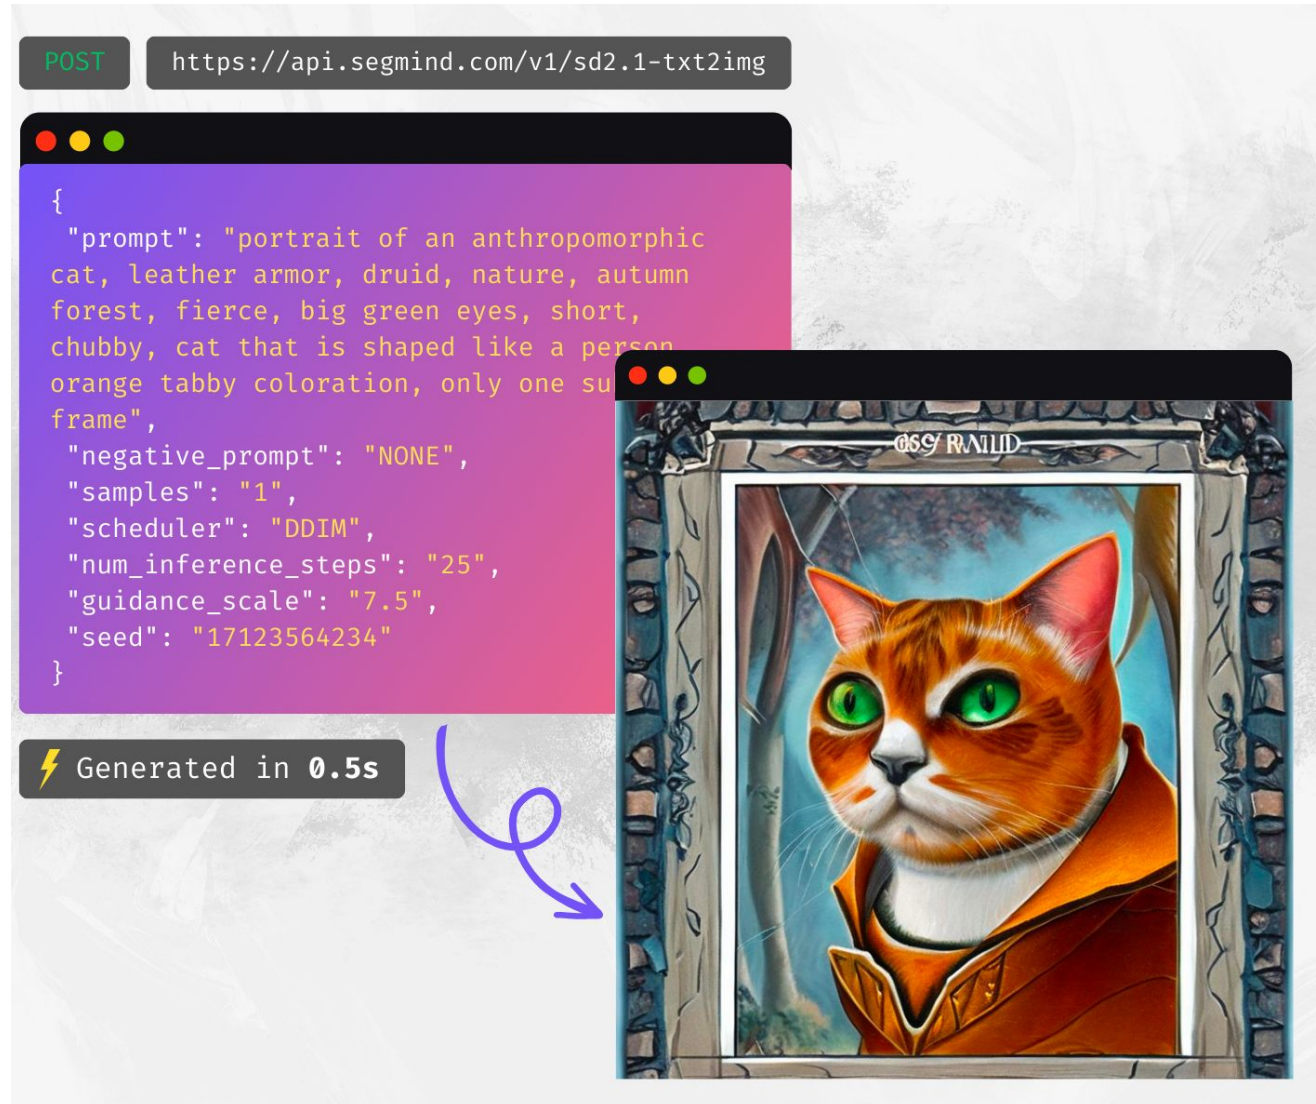
- Comment détermine un LLM

Text-to-image

Domain d'application:

- Image generation
- Image edition

<https://www.midjourney.com/>



# LLM: type de modèles

- Comment détermine un LLM

Text-to-3D  
(video)

Domain d'application:

- Video generation
- Video edition
- Jeux Vidéo



# LLM: type de modèles



GNOMON<sup>®</sup>  
DIGITAL

- Comment détermine un LLM

Text-to-Task (Tool GPT)

Domain d'application:

- Assistants virtuelle
- Automation
- LLM agent



[Ray-BAN META](#)



# LLM: type de modèles



GNOMON<sup>®</sup>  
DIGITAL

## Large Language Models

BCV

Application Layer

### Copywriting



### Coding



### Dev Tools



### Chat / Comms



### BizOps



Infrastructure Layer


### Model Creation



### Hardware



### Fine Tuning



### Inference



# Qu'est-ce que c'est Prompts et Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

 Le **prompt engineering** dans le contexte des modèles de langage (LLM, pour *Large Language Models*) désigne l'art et la science de concevoir des instructions textuelles (ou *prompts*) afin d'obtenir des résultats spécifiques ou optimaux d'un modèle de langage.

## Définition

Le prompt engineering consiste à rédiger, structurer ou formuler des *prompts* pour orienter le comportement d'un LLM et maximiser sa capacité à produire des réponses pertinentes, précises ou créatives. Cela peut inclure l'utilisation de formats particuliers, de questions claires, de contextes détaillés ou de contraintes explicites.

## Objectifs principaux

- **Optimisation des réponses** : Obtenir des résultats qui correspondent aux attentes en termes de contenu, style ou précision.
- **Réduction des ambiguïtés** : Limiter les interprétations multiples pour guider le modèle vers une réponse unique et claire.
- **Exploration créative** : Stimuler des réponses originales en utilisant des formulations non conventionnelles.

Généré par ChatGPT

## Techniques courantes

1. **Précision des consignes** : Formuler des demandes explicites avec des détails clairs (ex. "Expliquez en trois points la théorie de la relativité").
2. **Contexte riche** : Fournir un contexte pertinent pour améliorer la compréhension (ex. "Imaginez que vous êtes un scientifique du 19e siècle...").
3. **Format structuré** : Exiger des réponses sous forme de liste, tableau ou paragraphe (ex. "Donnez un résumé sous forme de liste à puces").
4. **Chain-of-thought prompting** : Encourager le raisonnement pas à pas pour résoudre des problèmes complexes.
5. **Exemples dans le prompt** : Inclure des exemples explicites pour clarifier les attentes (ex. "Voici un exemple de réponse idéale...").



# Qu'est-ce que c'est Prompts et Prompt Engineering




Il y a 3 façon d'entraîner un modèle LLM, chacun a besoin une façon de 'prompting' différente:

- Modèle LLM générique (ou brut)
- Optimisé par les instructions
- Optimisé par un dialogue

# Modèle LLM générique (ou brut)



GNOMON<sup>®</sup>  
DIGITAL




Token is a part of a word,  
the atomic unit that LLMs work in.

# Optimisé par les instructions



Entraîner à prédire une réponse à partir des instructions données dans les texts entrées



# Optimisé par une dialogue



GNOMON<sup>®</sup>  
DIGITAL

Entraîné à engager un dialogue en prédisant la prochaine réponse.



## Prompt examples

```
[User] Is the comment "do you like the  
weather?" ok or toxic?  
[Bot] ok.  
[User] can you briefly say why?  
[Bot]
```


## Model Output

```
It's just a question about the weather, people  
are not usually upset by that.
```

# Les utilisations LLM



GNOMON<sup>®</sup>  
DIGITAL



LLM peut être utilisé dans beaucoup de contextes différentes.




# Fin-tuning in LLM



GNOMON<sup>®</sup>  
DIGITAL

Le processus d'adaptation d'un modèle à un nouveau domaine ou à un ensemble de cas d'utilisation spécifiques en entraînant le modèle sur de nouvelles données. Par exemple, nous pouvons collecter des données d'entraînement et "ajuster" le modèle LLM spécifiquement pour le domaine juridique ou médical.




# Fin-tuning in LLM



GNOMON<sup>®</sup>  
DIGITAL

Le processus d'adaptation d'un modèle à un nouveau domaine ou à un ensemble de cas d'utilisation spécifiques en entraînant le modèle sur de nouvelles données. Par exemple, nous pouvons collecter des données d'entraînement et "ajuster" le modèle LLM spécifiquement pour le domaine juridique ou médical.



# Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

The **response** from  
the model



The **answers** you get depend on the  
**questions** you ask.



The **prompt** you  
designed

# Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

On teste les cas suivants:

- Who is Linus Torvalds?
- Summarize some important dates in the life of Linus Torvalds
- Write me a tweet to celebrate Linus Torvald's birthday
- Tester toi-même un prompt



## Elements of a Good Prompt

Act as a helpful tutor who breaks down complex subjects into easy explanations.

I want you to explain the process of photosynthesis to a

14 year old student, to assist with biology exam preparations.

Your answer should be 300 words, written in a tone that's friendly and educational.

**Persona:** Ask the tool to take a role

**Objective:** What do you want the AI to do

**Audience:** Specify who it's for

**Context:** What does the tool need to know

**Boundaries:** Set your own direction & limitation

*Tip 1*

### Give Clear Instructions

Use commands that instruct the AI tool on what you want to generate, such as 'explain', 'translate', 'summarize' or 'compare'.

*Tip 2*

### Provide Context

Adding context and background information can help the tool to understand the task better. For example, mention the project type such as 'short story', 'report' or 'outline'

*Tip 3*

### Iterate & Experiment

Try different instructions and techniques if you don't get the results you want. Prompting can be like an experiment that may require several rounds of iterations!



## Ambiguïté et Nuance:

- **Prompt A (Ambigu):** "Écris une histoire."
- **Prompt B (Précis):** "Écris une histoire courte de science-fiction sur un robot qui découvre qu'il est capable de ressentir des émotions, en utilisant un vocabulaire riche et des descriptions imagées."

## Contraintes Créatives

- **Prompt 5 (Libre):** "Écris un poème."
- **Prompt 6 (Contraint):** "Écrit un sonnet en alexandrins sur le thème de la solitude, en utilisant au moins trois fois le mot 'silence'."

## Instructions Complexes

- **Prompt A (Simple):** "Traduis 'Hello, how are you?' en français."
- **Prompt B (Complexe):** "Traduis 'Hello, how are you?' en français, puis réécris la traduction en utilisant un langage formel, puis informel, et enfin en argot parisien. Indique pour chaque version le contexte approprié."

## Raisonnement et Connaissance


- **Prompt A (Factuel):** "Qui a peint la Joconde?"
- **Prompt B (Raisonnement):** "Si tous les chats sont des mammifères, et que Minou est un chat, est-ce que Minou est un mammifère ? Explique ton raisonnement."

## Sensibilité au contexte

- **Prompt 9 (Sans contexte):** "Explique le concept de la relativité."
- **Prompt 10 (Avec contexte):** "Imagine que tu expliques le concept de la relativité à un enfant de 8 ans. Utilise des analogies simples et évite le jargon scientifique."

### Tips:

- By nice
- Step by step
- Ask LLM to write your prompt



# POWER OF PROMPT ENGINEERING

## ZERO-SHOT INFERENCE

A MODEL THAT CAN  
ANSWER QUESTIONS  
OR PERFORM TASKS  
WITHOUT ANY  
SPECIFIC TRAINING ON  
THOSE PARTICULAR  
PROMPTS.

## ONE-SHOT INFERENCE

ONE-SHOT INFERENCE  
TAKES ONE EXAMPLE. THE  
MODEL CAN GRASP THE  
ESSENCE OF THE TASK AND  
GENERATE THE DESIRED  
OUTPUT.

## FEW-SHOT INFERENCE

FEW-SHOT INFERENCE  
ALLOWS YOU TO PROVIDE  
A SMALL NUMBER OF  
EXAMPLES TO GUIDE THE  
MODEL'S BEHAVIOR. IT'S  
LIKE GIVING THE MODEL A  
MINI-TRAINING SESSION  
WITH JUST A HANDFUL  
OF PROMPTS.

<https://www.promptingguide.ai/>

# Zero-shot prompting



On donne des instructions à LLM sans aucun exemples

EX:

Quelle est la sentiment de la phrase suivant: 'I am very happy today'? Positive, négative ou neutre.

<https://www.promptingguide.ai/>

# One-shot prompting



GNOMON<sup>®</sup>  
DIGITAL

On donne des instructions à LLM avec un exemple

EX:

Quelle est la sentiment de la phrase suivant: 'I am very happy today'? Positive, négative ou neutre.

Example:

Input: "The book is very interesting"

Output: "Positive"

<https://www.promptingguide.ai/>



# Prompt Engineering



GNOMON<sup>®</sup>  
DIGITAL

On teste?

- Gemini: [https://aistudio.google.com/prompts/new\\_chat](https://aistudio.google.com/prompts/new_chat)
- ChatGPT: <https://chatgpt.com/>

<https://www.promptingguide.ai/>

# Few-shot prompting



GNOMON<sup>®</sup>  
DIGITAL

On donne des instructions à LLM avec un exemple

EX:

Quelle est la sentiment de la phrase suivant: 'I am very happy today'? Positive, négative ou neutre.

Example:

Input: "The book is very interesting"

Output: "Positive"

Input: "I feel so depressed today"

Output: "Négative"

Input: "My name is John"

Output: "Neutre"

<https://www.promptingguide.ai/>

# Few-shot prompting



[https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex\\_genai/labs/prompt\\_engineering.ipynb](https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex_genai/labs/prompt_engineering.ipynb)

[https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex\\_genai/labs/gemini\\_for\\_multimodal\\_prompting.ipynb](https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex_genai/labs/gemini_for_multimodal_prompting.ipynb)

# Chain of Thought



GNOMON<sup>®</sup>  
DIGITAL

Magique prompt: “Let’s think this through step by step”

## Standard Prompting

### Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The answer is 27. ❌

## Chain of Thought Prompting

### Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output


A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9. ✅

# Self Refine



GNOMON<sup>®</sup>  
DIGITAL

## Self-Refine



Learn Prompting

### SELF-REFINE: Iterative Refinement with Self-Feedback

Aman Madaan<sup>1</sup>, Niket Tandon<sup>2</sup>, Prakhar Gupta<sup>1</sup>, Skyler Hallinan<sup>3</sup>, Luyu Gao<sup>1</sup>, Sarah Wiegrefe<sup>2</sup>, Uri Alon<sup>1</sup>, Nouha Dziri<sup>2</sup>, Shrimai Prabhumoye<sup>4</sup>, Yiming Yang<sup>1</sup>, Shashank Gupta<sup>2</sup>, Bodhisattwa Prasad Majumder<sup>5</sup>, Katherine Hermann<sup>6</sup>, Sean Welleck<sup>2,3</sup>, Amir Yazdanbakhsh<sup>6</sup>, Peter Clark<sup>2</sup>

<sup>1</sup>Language Technologies Institute, Carnegie Mellon University

<sup>2</sup>Allen Institute for Artificial Intelligence

<sup>3</sup>University of Washington <sup>4</sup>NVIDIA <sup>5</sup>UC San Diego <sup>6</sup>Google Research, Brain Team  
amadaan@cs.cmu.edu, nikett@allenai.org

# Self Consistency



GNOMON<sup>®</sup>  
DIGITAL

## Self-consistency

**Q:** If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

**A:** There are 3 cars in the parking lot already. 2 more arrive. Now there are  $3 + 2 = 5$  cars. The answer is 5.

...

**Q:** Janet's ducks lay 16 eggs per day. She eats three for breakfast every morning and bakes muffins for her friends every day with four. She sells the remainder for \$2 per egg. How much does she make every day?

**A:**

Language  
model

Sample a diverse set of  
reasoning paths

She has  $16 - 3 - 4 = 9$  eggs left. So she makes  $\$2 * 9 = \$18$  per day.

**The answer is \$18.**

This means she she sells the remainder for  $\$2 * (16 - 4 - 3) = \$26$  per day.

**The answer is \$26.**

She eats 3 for breakfast, so she has  $16 - 3 = 13$  left. Then she bakes muffins, so she has  $13 - 4 = 9$  eggs left. So she has  $9 \text{ eggs} * \$2 = \$18$ .

**The answer is \$18.**

Marginalize out reasoning paths  
to aggregate final answers

**The answer is \$18.**



# Self Consistency



GNOMON<sup>®</sup>  
DIGITAL

## Stage 1: Decompose Question into Subquestions

Q: It takes Amy 4 minutes to climb to the top of a slide. It takes her 1 minute to slide down. The water slide closes in 15 minutes. How many times can she slide before it closes?

Language Model

A: To solve "How many times can she slide before it closes?", we need to first solve: "How long does each trip take?"

## Stage 2: Sequentially Solve Subquestions




Figure 1: Least-to-most prompting solving a math word problem in two stages: (1) query the language model to decompose the problem into subproblems; (2) query the language model to sequentially solve the subproblems. The answer to the second subproblem is built on the answer to the first subproblem. The demonstration examples for each stage's prompt are omitted in this illustration.

# Tree of Thoughts



## Tree of Thoughts: Deliberate Problem Solving with Large Language Models

Shunyu Yao  
Princeton University

Dian Yu  
Google DeepMind

Jeffrey Zhao  
Google DeepMind

Izhak Shafran  
Google DeepMind

Thomas L. Griffiths  
Princeton University

Yuan Cao  
Google DeepMind

Karthik Narasimhan  
Princeton University




Figure 1: Schematic illustrating various approaches to problem solving with LLMs. Each rectangle box represents a *thought*, which is a coherent language sequence that serves as an intermediate step toward problem solving. See concrete examples of how thoughts are generated, evaluated, and searched in Figures 2,4,6.



La limite de LLM:

- Pas d'accès aux informations à jour

Quelle est le resultat d'élection présidentielle 2024

Cancel

Send



Je n'ai pas accès aux résultats de l'élection présidentielle de 2024 sans effectuer une recherche en direct, et je ne dispose pas de données en temps réel intégrées à ma mémoire. Si vous souhaitez connaître les résultats actualisés, je peux effectuer une recherche pour vous. 😊

# RAG



GNOMON<sup>®</sup>  
DIGITAL



Searched 2 sites ^

Q "résultat élection présidentielle 2024"  
bing.com

 Le Monde.fr — Présidentielle américaine 2024 : retrouvez les dernières informations après l'élection ...  
lemonde.fr

 CNEWS — Présidentielle américaine 2024 : suivez les résultats de l'élection en temps réel sur notre...  
cnews.fr

Lors de l'élection présidentielle américaine de 2024, Donald Trump a été déclaré vainqueur face à Kamala Harris. Trump a obtenu 312 votes des grands électeurs, dépassant largement la majorité requise de 270, tandis que Harris en a obtenu 226. Ce résultat marque le retour de Trump à la Maison-Blanche pour un second mandat après une campagne acharnée. L'investiture officielle est prévue pour le 20 janvier 2025 LE MONDE.FR CNEWS .


< 2/2 > 🔊 📄 👍 🗨️ ↺

# RAG



GNOMON<sup>®</sup>  
DIGITAL

## RAG Architecture Model





# RAG: recherche sémantique

**La notion d'adéquation sémantique entre deux entités X et Y s'applique à de nombreux contextes :**

- Le sens de Y est-il similaire à celui de X ?
- X et Y représentent-ils la même chose présentée différemment ?
- Y est-il une bonne recommandation ou une réponse appropriée pour X ?

## Lexical search

daniel radcliffe

"daniel" + "radcliffe"




## Semantic search

how rich is the actor who played harry potter

"Net worth"

"Daniel Radcliffe"





# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL

## Lexical Search

## Semantic Search

### Definition

Searches based on exact keyword match

Understands query intent and context

### Matching Criteria

Matches keywords or phrases literally

Considers synonyms, context, and relationships between words

### Natural Language Processing

Requires minimal NLP capabilities

Requires advanced NLP for understanding context


### Example

Searching for "Amazon river" returns info about amazon river, amazon.com etc.

Searching for "Amazon river" returns information about the Amazon river and not the company Amazon.




## Entity Resolution



# RAG: researcher semantic





## Process of Facial Recognition Payment




# RAG: recherche sémantique




GNOMON<sup>®</sup>  
DIGITAL



L2 distance



cosine similarity



inner product

Calculating vector similarity



2 dimensions x  
1M items  
=  $O(1M \times 2)$

# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL


Data




Embeddings



Search engine



Index







# NLP



GNOMON<sup>®</sup>  
DIGITAL




word2vec



# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL




Very similar to word2vec

# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL

## Step 1: Tokenize



Normalization

Pre-tokenization

Model

Postprocessor

# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL

Step 2: Learn embedding vectors

[100, 2031, 3, 18, 10, 8, ..]

|     |     |     |     |     |
|-----|-----|-----|-----|-----|
| ↓   | ↓   | ↓   | ↓   | ... |
| 0.3 | 0.6 | 0.1 | 1.9 |     |
| 1.2 | 2.1 | 1.9 | 1.5 |     |
| 1.6 | 2.2 | 2.1 | 0.3 |     |
| 0.9 | 1.3 | 1.8 | 2.2 |     |
| 0.3 | 1.5 | 0.8 | 1.8 |     |
| 0.5 | 1.0 | 0.9 | 1.1 |     |
| 1.6 | 1.4 | 3.2 | 0.4 |     |
| 2.3 | 0.6 | 1.1 | 1.6 |     |



Step 2: Learn embedding vectors

## Different types of embeddings

Type 1

Non-Contextual  
Embeddings

Factorization Based

- Word2Vec
- Gloves
- Swivel

Type 2

Contextual  
Embeddings

Language Model Based

- Bert
- PaLM
- GPT

# RAG: recherche semantic




Step 2: Learn embedding vectors

## Different types of embeddings


Type 1

Non-Contextual  
Embeddings



Type 2

Contextual  
Embeddings



# RAG: recherche semantic




Step 2: Learn embedding vectors

## Different types of embeddings


Type 1

Non-Contextual  
Embeddings



Type 2

Contextual  
Embeddings






# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL


## MATRIX FACTORIZATION




# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL



# RAG: recherche semantic



# RAG: recherche semantic




GNOMON<sup>®</sup>  
DIGITAL

Step 2: Learn embedding vectors

## Different types of embeddings


Type 1

Non-Contextual  
Embeddings



Type 2

Contextual  
Embeddings




# RAG: researcher semantic



GNOMON<sup>®</sup>  
DIGITAL

## Contextual: Recurrent Network Language Models




- Sequential computation
- Context only from one side
- Can be stacked

# RAG: researcher semantic



GNOMON<sup>®</sup>  
DIGITAL

## Contextual: Transformer Language Models







# RAG: recherche sémantique




GNOMON<sup>®</sup>  
DIGITAL



L2 distance



cosine similarity



inner product

2 dimensions x  
1M items  
=  $O(1M \times 2)$


Calculating vector similarity

L2 distance (Euclidean) distance:

## K-Nearest Neighbors: Distance Metric


L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$






# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL

Cos similarity:

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$



# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL

Dot product:


$$\mathbf{a} = \langle a_1, a_2, a_3 \rangle \quad \mathbf{b} = \langle b_1, b_2, b_3 \rangle$$

$$\mathbf{a} \cdot \mathbf{b} = a_1 b_1 + a_2 b_2 + a_3 b_3$$


# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL



$$\mathbf{a} \cdot \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos \alpha$$




$$\text{sim}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{||\mathbf{a}|| \cdot ||\mathbf{b}||}$$

# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL

## Similarity Distance Measure





# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL


[https://github.com/GoogleCloudPlatform/ai-ml-immersion/blob/master/notebooks/vertex\\_genai/labs/semantic\\_matching\\_with\\_gemini.ipynb](https://github.com/GoogleCloudPlatform/ai-ml-immersion/blob/master/notebooks/vertex_genai/labs/semantic_matching_with_gemini.ipynb)

# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL

## Vector search data workflow



# RAG: recherche semantic



GNOMON<sup>®</sup>  
DIGITAL


Recherche par vectorielle est splitté en deux parties:

- Offline -> Calcule lourde /lente:
  - Ajouter un vector index (keys= vectors)
  - Sauve tous les embedding vecteurs dans ces index
- Online -> Calcule efficace / Rapid
  - Algorithm search dans vectorDB

# RAG: researcher semantic



GNOMON<sup>®</sup>  
DIGITAL



# RAG: recherche sémantique



GNOMON<sup>®</sup>  
DIGITAL

[https://github.com/GoogleCloudPlatform/ai-ml-immersion/blob/master/notebooks/vertex\\_genai/solutions/semantic\\_search\\_with\\_vector\\_search.ipynb](https://github.com/GoogleCloudPlatform/ai-ml-immersion/blob/master/notebooks/vertex_genai/solutions/semantic_search_with_vector_search.ipynb)



# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL





# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

Untitled prompt 

← Compare   > Get code   ⋮

---

^ System Instructions

Optional tone and style instructions for the model

---

## What will you build?

Push Gemini to the limits of what AI can do

 **Trip recommendations**  
Convert unorganized text into structured tables.

 **Unit Testing**  
Add unit tests for a Python function.

 **Modify writing style**  
Change the tone and writing style of a blurb.

**Run settings** Reset

 **Model**  
Gemini 1.5 Flash 

 **Token count**  
0 / 1000 000

 **Temperature**  
 1

^ **Tools**

**Structured output**  
 [Edit schema](#)

**Code execution**  


**Function calling**  
 [Edit functions](#)



Tasks

- Classification
- Summarization
- Q&A




[https://aistudio.google.com/app/prompts/new\\_chat](https://aistudio.google.com/app/prompts/new_chat)

# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

## Tool LLM



## TOOLLLM: FACILITATING LARGE LANGUAGE MODELS TO MASTER 16000+ REAL-WORLD APIS

Yujia Qin<sup>1\*</sup>, Shihao Liang<sup>1\*</sup>, Yining Ye<sup>1</sup>, Kunlun Zhu<sup>1</sup>, Lan Yan<sup>1</sup>, Yaxi Lu<sup>1</sup>, Yankai Lin<sup>3†</sup>, Xin Cong<sup>1</sup>, Xiangru Tang<sup>4</sup>, Bill Qian<sup>4</sup>, Sihan Zhao<sup>1</sup>, Lauren Hong<sup>1</sup>, Runchu Tian<sup>1</sup>, Ruobing Xie<sup>5</sup>, Jie Zhou<sup>5</sup>, Mark Gerstein<sup>4</sup>, Dahai Li<sup>2,6</sup>, Zhiyuan Liu<sup>1†</sup>, Maosong Sun<sup>1†</sup>

<sup>1</sup>Tsinghua University <sup>2</sup>ModelBest Inc. <sup>3</sup>Renmin University of China


<sup>4</sup>Yale University <sup>5</sup>WeChat AI, Tencent Inc. <sup>6</sup>Zhihu Inc.

yujiaqin16@gmail.com

# Introduction à LangChain et LlamaIndex



LLM Agent



# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

LLM Agent

## Exploring Large Language Model based Intelligent Agents: Definitions, Methods, and Prospects

Yuheng Cheng<sup>1</sup> <sup>1</sup> Ceyao Zhang<sup>1</sup> <sup>1</sup> Zhengwen Zhang<sup>1</sup>

Xiangrui Meng<sup>1</sup> Sirui Hong<sup>2</sup> Wenhao Li<sup>1</sup> Zihao Wang<sup>3</sup> Zekai Wang<sup>4</sup>

Feng Yin<sup>1</sup> Junhua Zhao<sup>1</sup> Xiuqiang He<sup>5</sup>

<sup>1</sup>The Chinese University of Hong Kong, Shenzhen


<sup>2</sup>DeepWisdom <sup>3</sup>Peking University <sup>4</sup>Yantu.ai <sup>5</sup>FiT, Tencent

These authors contributed to the work equally and should be regarded as co-first authors. Corresponding author

# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL



Usages possibles:

- Question & Answering over private corpus
- Large document Summarization
- Personal assistants & Chatbots
- Querying tabular data
- Code understanding
- API interactions
- Extraction
- Evaluation



# Introduction à LangChain et LlamaIndex



## LlamaIndex vs. LangChain




| Features             |  LlamaIndex |  LangChain |
|----------------------|----------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| <b>Primary focus</b> | Intelligent search, data indexing and retrieval                                              | Building a wide range of Gen AI applications                                                  |
| <b>Data handling</b> | Ingesting, structuring, and accessing private or domain-specific data                        | Loading, processing, and indexing data for various uses                                       |
| <b>Customization</b> | Offers tools for integrating private data into LLMs                                          | Highly customizable, users can chain multiple components                                      |
| <b>Flexibility</b>   | Specialized for efficient and fast search                                                    | General-purpose framework with flexibility in application behavior                            |
| <b>Deployment</b>    | Ideal for proprietary or specialized data                                                    | Facilitates the deployment of bespoke NLP applications                                        |
| <b>Use cases</b>     | Best for applications that require quick data lookup and retrieval                           | Suitable for applications that require complex interactions like chatbots, GQA, summarization |



# Introduction à LangChain



GNOMON<sup>®</sup>  
DIGITAL




LangChain  
Framework




LangChain

# Introduction à LangChain et LlamaIndex



# Introduction à LangChain et LlamaIndex



# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

- Model IO focuse sur les fonctionnalités basique de LLM (text entrée/sortie)
- Avantage LangChain:
  - Une seul interface aux tous les providers LLMs (Google, OpenAI et Hugging Face etc)
  - Text mode et Chat mode
  - Pas d'adaptation à faire si on change le modèle LLM

[https://python.langchain.com/docs/tutorials/llm\\_chain/](https://python.langchain.com/docs/tutorials/llm_chain/)

# Introduction à LangChain et LlamaIndex



- Data Connection
  - Possibilité de connecter LLM avec des données externes (RAG)
- Avantages:
  - Integration functions disponibles aux data sources populaires (CSVs, PDFs, AWS, Google Cloud, etc)
  - Possible d'utiliser/changer la solution VectorDB

[https://python.langchain.com/docs/integrations/document\\_loaders/](https://python.langchain.com/docs/integrations/document_loaders/)

# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

```
from langchain.vectorstores import Chroma
from langchain.embeddings import VertexAIEmbeddings
from langchain_community.document_loaders import PyPDFDirectoryLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter

documents = PyPDFDirectoryLoader("pdfs/").load()

text_splitter = RecursiveCharacterTextSplitter(chunk_size=800, chunk_overlap=400)

document_chunks = text_splitter.split_documents(documents)

embedding = VertexAIEmbeddings(model_name="textembedding-gecko@001")

db = Chroma.from_documents(document_chunks, embedding,
persist_directory="./vectorstore")
```

[https://python.langchain.com/docs/integrations/document\\_loaders/](https://python.langchain.com/docs/integrations/document_loaders/)



# Introduction à LangChain et LlamaIndex



- Chains :
  - Les sorties d'un LLM peuvent être les entrées d'un autre LLM (Appel APIs LLM en chain)
- Avantages:
  - Facilite l'utilisation d'appelle en chains


# Introduction à LangChain et LlamaIndex



## stanford-futuredata/ FrugalGPT

FrugalGPT: better quality and lower cost for LLM applications

2 Contributors  
3 Issues  
187 Stars  
22 Forks




(a) Learned FrugalGPT strategy




(b) A query and response example

| Approach  | Accuracy | Cost (\$) |
|-----------|----------|-----------|
| GPT-4     | 0.857    | 33.1      |
| FrugalGPT | 0.872    | 6.5       |

(c) Overall performance and cost



(b) Proposed FrugalGPT




(c) Performance and cost tradeoffs

Figure 3: A case study of FrugalGPT on the HEADLINES dataset. (a) The cascade strategy that FrugalGPT learned on this dataset with overall budget \$6.5, one fifth of GPT-4's cost. FrugalGPT avoids querying GPT-4 as long as GPT-J and J1-L produce high-quality answers. (b) Sometimes GPT-4 makes a mistake, but FrugalGPT learns to use the correct answers by J-1 and GPT-J. (c) Overall, we observe that FrugalGPT reduces the cost by 80%, while improves the accuracy by 1.5% compared to GPT-4.

# Introduction à LangChain et LlamaIndex



LLM Agent



# Introduction à LangChain et LlamaIndex



GNOMON<sup>®</sup>  
DIGITAL

- Memory :
  - Laisser tes modèles d'y accès des contexts historiques des précédents conversations

```
from langchain.llms import VertexAI
from langchain.chains import ConversationChain
from langchain.memory import ConversationBufferMemory

llm = VertexAI(model_name="codechat-bison")

conversation = ConversationChain(
    llm=llm, verbose=True, memory=ConversationBufferMemory()
)
```



La limite de LLM:

- Pas d'accès aux informations à jour

Quelle est le resultat d'élection présidentielle 2024

Cancel

Send



Je n'ai pas accès aux résultats de l'élection présidentielle de 2024 sans effectuer une recherche en direct, et je ne dispose pas de données en temps réel intégrées à ma mémoire. Si vous souhaitez connaître les résultats actualisés, je peux effectuer une recherche pour vous. 😊

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL



Searched 2 sites ^

Q "résultat élection présidentielle 2024"  
bing.com

 Le Monde.fr — Présidentielle américaine 2024 : retrouvez les dernières informations après l'élection ...  
lemonde.fr

 CNEWS — Présidentielle américaine 2024 : suivez les résultats de l'élection en temps réel sur notre...  
cnews.fr

Lors de l'élection présidentielle américaine de 2024, Donald Trump a été déclaré vainqueur face à Kamala Harris. Trump a obtenu 312 votes des grands électeurs, dépassant largement la majorité requise de 270, tandis que Harris en a obtenu 226. Ce résultat marque le retour de Trump à la Maison-Blanche pour un second mandat après une campagne acharnée. L'investiture officielle est prévue pour le 20 janvier 2025 LE MONDE.FR CNEWS .

< 2/2 > 🔊 📄 👍 🗨️ ↺



# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL



Searched 2 sites ^

Q "résultat élection présidentielle 2024"  
bing.com

 Le Monde.fr — Présidentielle américaine 2024 : retrouvez les dernières informations après l'élection ...  
lemonde.fr

 CNEWS — Présidentielle américaine 2024 : suivez les résultats de l'élection en temps réel sur notre...  
cnews.fr

Lors de l'élection présidentielle américaine de 2024, Donald Trump a été déclaré vainqueur face à Kamala Harris. Trump a obtenu 312 votes des grands électeurs, dépassant largement la majorité requise de 270, tandis que Harris en a obtenu 226. Ce résultat marque le retour de Trump à la Maison-Blanche pour un second mandat après une campagne acharnée. L'investiture officielle est prévue pour le 20 janvier 2025 LE MONDE.FR CNEWS .

< 2/2 > 🔊 📄 👍 🗣️ ↺

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL



Who was the first person to walk on the moon?



**Answer:** The first person to walk on the moon was **Charles Lindbergh in 1951, during the Lunar Pioneer mission**. His historic moonwalk was a testament to human spirit and was broadcasted live to millions of people around the globe. ❌



**Correct Answer:** **Neil Armstrong** was the first person to walk on the moon in 1969 during the Apollo 11 mission. ✅

(a) Factuality Hallucination



Please summarize the following news article:



**Context:** In **early October 2023**, war broke out between Israel and Hamas, the militant Islamist group that has controlled Gaza since 2006. Hamas fighters fired rockets ... civilians and taking dozens of hostages.

**Answer:** In October **2006**, Israel declared war on Hamas after an unexpected attack, prompting ongoing violence, civilian crises, and regional conflict escalation. ❌

(b) Faithfulness Hallucination

Figure 1: An intuitive example of LLM hallucination.

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

## (Full) Fine-Tuning

Run fine-tuning to teach the LLM by baking knowledge into the weights of the network.

- Data preparation effort
- Expensive to do
- Online learning/ updating new data
- Still may not work

## Make Humans Check

Put a human in-between the output of the LLM and the user.

- Very latent and expensive
- Humans untrustworthy or inattentive

## Prompt Engineering

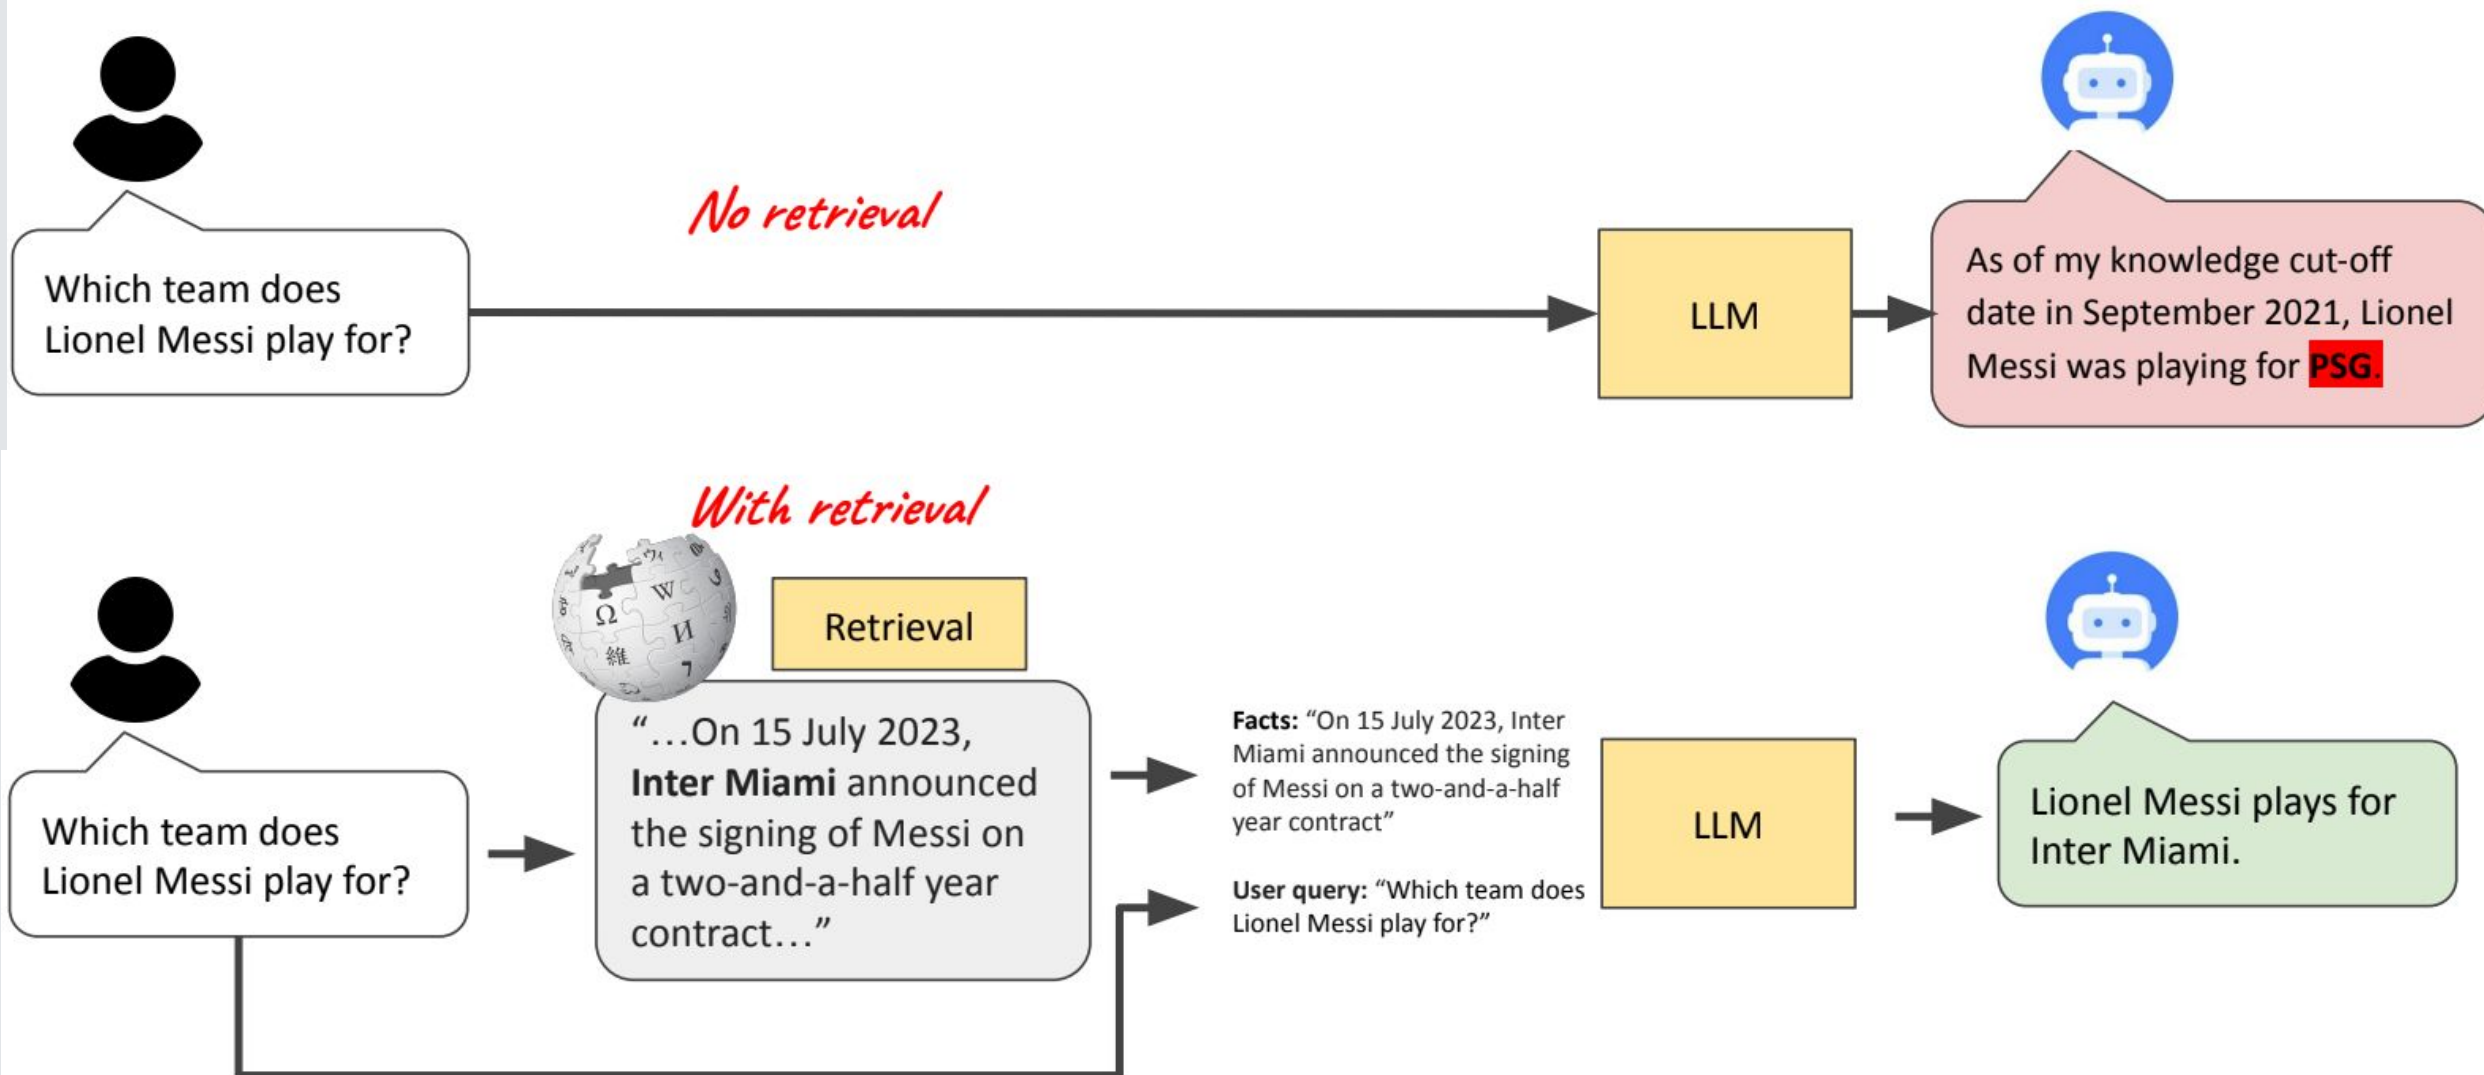
Induce the LLM by adding any known information to the prompt.

- Based on LLMs existing trained knowledge
- Token limit for LLMs
- Error prone: Retrieving the right context
- Trade off: performance, latency, cost

# Grounding with Public Knowledge




GNOMON<sup>®</sup>  
DIGITAL



# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

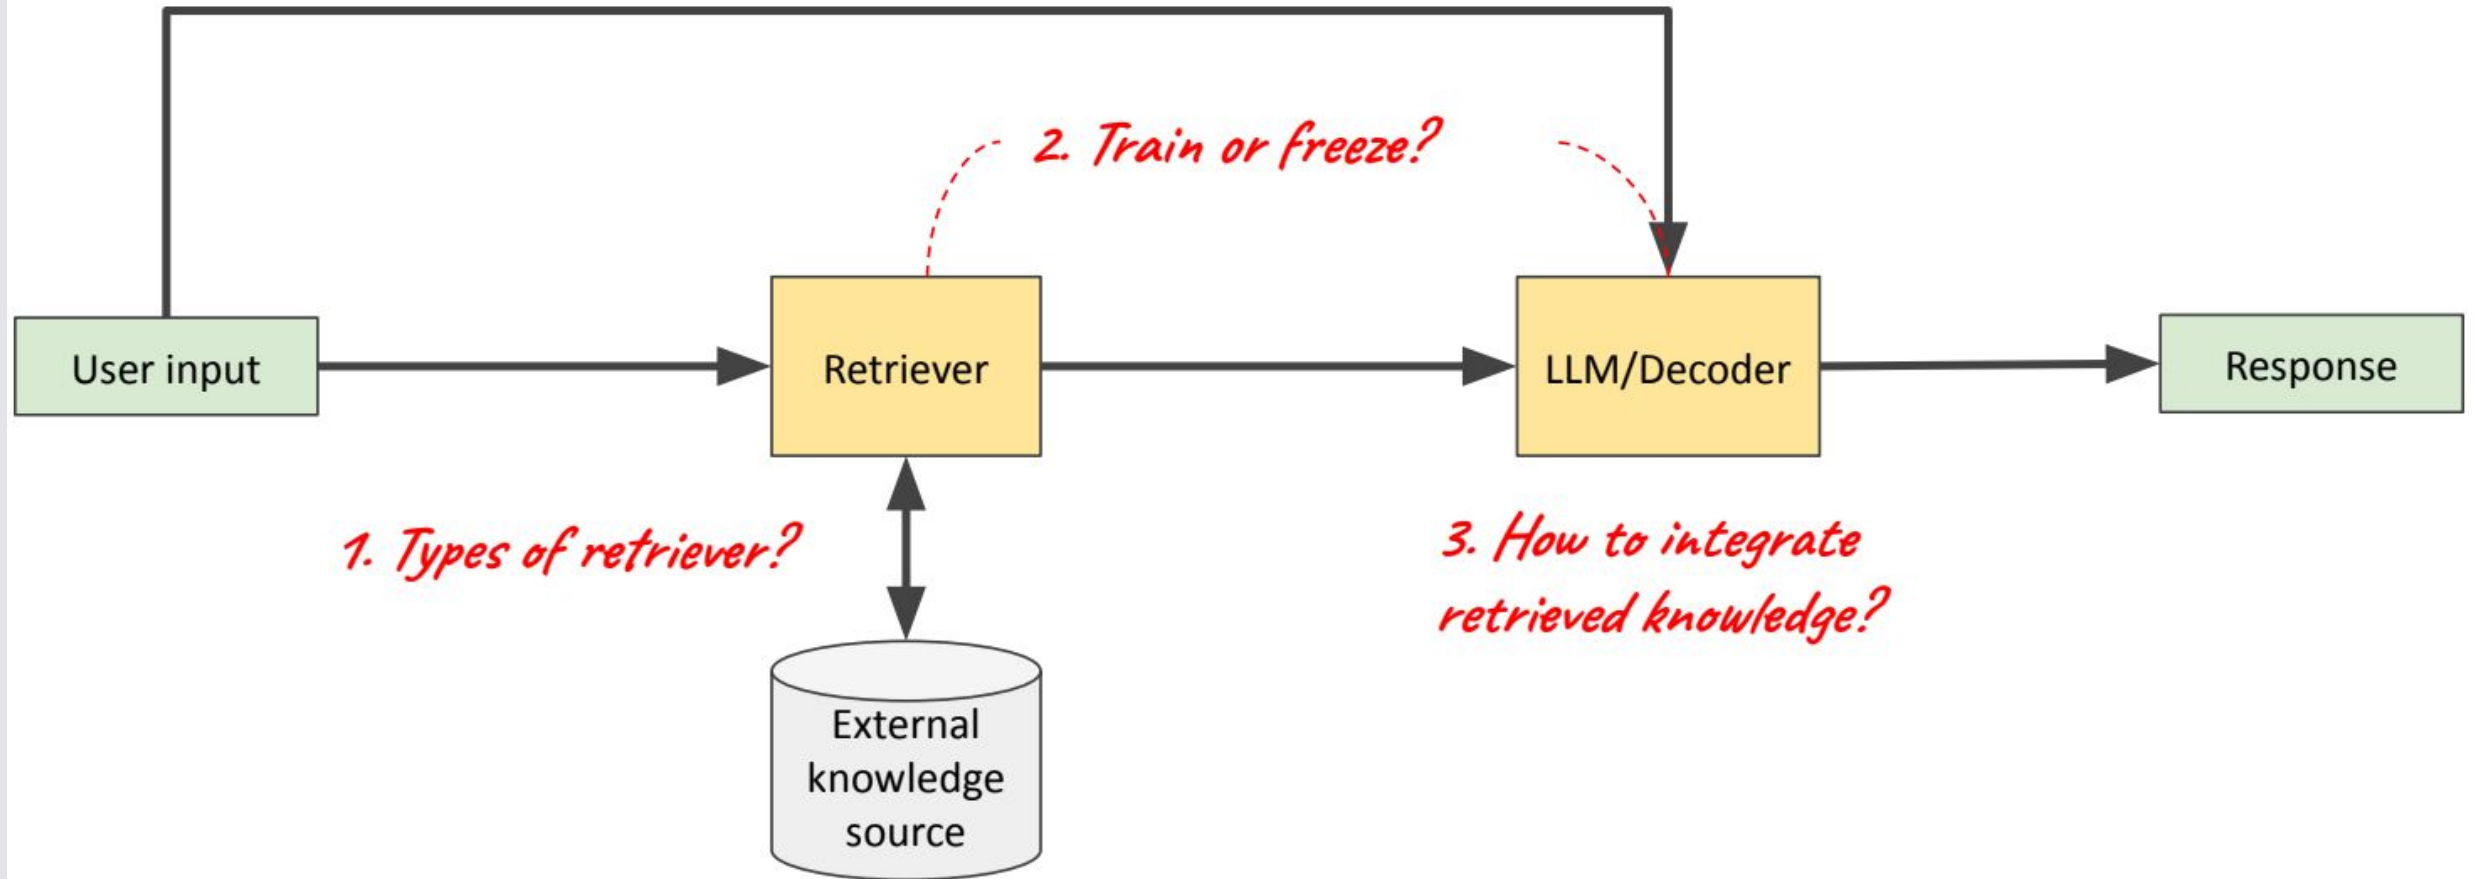




# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL






# Grounding with Public Knowledge



- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity
- Knowledge Graphs
  - Graph traversal (e.g. n-hop)
  - Autoregressive path finding
  - GraphRag
- Web:
  - Off-the-shelf search engine for retrieval




# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity
- Knowledge Graphs
  - Graph traversal (e.g. n-hop)
  - Au
  - Gr
- Web:
  - Of



$$w_{x,y} = tf_{x,y} \times \log \left( \frac{N}{df_x} \right)$$

**TF-IDF**

Term  $x$  within document  $y$

$tf_{x,y}$  = frequency of  $x$  in  $y$

$df_x$  = number of documents containing  $x$

$N$  = total number of documents

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

| Word    | TF  |     | IDF               | TF*IDF |       |
|---------|-----|-----|-------------------|--------|-------|
|         | A   | B   |                   | A      | B     |
| The     | 1/7 | 1/7 | $\log(2/2) = 0$   | 0      | 0     |
| Car     | 1/7 | 0   | $\log(2/1) = 0.3$ | 0.043  | 0     |
| Truck   | 0   | 1/7 | $\log(2/1) = 0.3$ | 0      | 0.043 |
| Is      | 1/7 | 1/7 | $\log(2/2) = 0$   | 0      | 0     |
| Driven  | 1/7 | 1/7 | $\log(2/2) = 0$   | 0      | 0     |
| On      | 1/7 | 1/7 | $\log(2/2) = 0$   | 0      | 0     |
| The     | 1/7 | 1/7 | $\log(2/2) = 0$   | 0      | 0     |
| Road    | 1/7 | 0   | $\log(2/1) = 0.3$ | 0.043  | 0     |
| Highway | 0   | 1/7 | $\log(2/1) = 0.3$ | 0      | 0.043 |

$$w_{x,y} = \text{tf}_{x,y} \times \log\left(\frac{N}{\text{df}_x}\right)$$

**TF-IDF**

Term  $x$  within document  $y$

$\text{tf}_{x,y}$  = frequency of  $x$  in  $y$

$\text{df}_x$  = number of documents containing  $x$




$N$  = total number of documents

# Grounding with Public Knowledge




- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity


- Knowledge Graphs



L2 distance



cosine similarity



inner product

Calculating vector similarity




2 dimensions x  
1M items  
=  $O(1M \times 2)$

# Grounding with Public Knowledge



- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity

- Knowledge Graphs






We can train the retriever

# Grounding with Public Knowledge



- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity

- Knowledge Graphs



We can train the retriever




# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

**Train the retriever to find documents that are most likely to reduce the perplexity of LLM**



$$\text{score}(v_1, v_2) = \frac{1}{1 + e^{-f_{\theta}(v_1, v_2)}}$$


neural network learned from data

$(v_1, v_2, \text{label})$


# Grounding with Public Knowledge



- Documents non structur  :
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity
- Knowledge Graphs
  - Graph traversal (e.g. n-hop)
  - Autoregressive path finding
  - GraphRag
- Web:
  - Off-the-shelf search engine for retrieval



# Grounding with Public Knowledge




A knowledge graph is made up of three main components: nodes, edges, and labels. Any object, place, or person can be a node. An edge defines the relationship between the nodes. For example, a node could be a client, like IBM, and an agency like, Ogilvy.

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL

## LONG-TAIL KEYWORDS



Long-tail keywords are search queries that get a small number of searches per month. They tend to be longer and more specific than their “head” counterparts and, therefore, often have a higher conversion rate.

# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL



google



<https://www.lovesdata.com/blog/2016/google-knowledge-graph>

[All](#) [Maps](#) [Images](#) [Videos](#) [News](#) [More ▾](#) [Search tools](#)

<https://support.google.com/knowledgepanel/answer/9787176?hl=en>

About 12,270,000,000 results (0.66 seconds)

## Google

<https://www.google.com.au/> ▾

Offers the choice of searching the whole web or web pages from Australia. Also advanced search, image and groups search, news and directory from the Open ...

### Google Maps

Find local businesses, view maps and get driving directions in ...

### Images

Google Images. The most comprehensive image search ...

### Google News

Aggregated headlines and search engine for many news services ...

### Google AdWords

Google AdWords lets you manage your campaign by yourself, or ...

[More results from google.com.au »](#)

## In the news



### A new look for Google Play family of apps

Official Android Blog - 2 days ago

Whether you like watching Despicable Me on Google Play Movies & TV, streaming "Sorry" ...

Here's How Google Makes Sure It (Almost) Never Goes Down

WIRED - 8 hours ago

Google To Donate IBM-Based Server Designs To Open Compute Foundation

Fortune - 2 hours ago

[More news for google](#)

## Google

Technology company



Google is an American multinational technology company specializing in Internet-related services and products. These include online advertising technologies, search, cloud computing, and software. [Wikipedia](#)

**CEO:** [Sundar Pichai](#)

**Founded:** September 4, 1998, Menlo Park, California, United States

**Headquarters:** Mountain View, California, United States

**Subsidiaries:** YouTube, ITA Software, AdMob, Google Japan, more

**Founders:** [Larry Page](#), [Sergey Brin](#)

**Parent organization:** [Alphabet Inc.](#)

## Profiles



Facebook



YouTube



Twitter



LinkedIn



Google+

[Feedback](#)

# Grounding with Public Knowledge



How to query a knowledge graph

Neo4J

| Cypher                                 | SPARQL                                                        |
|----------------------------------------|---------------------------------------------------------------|
| CREATE                                 | INSERT                                                        |
| RETURN                                 | SELECT                                                        |
| WITH                                   | SELECT in a sub query                                         |
| MATCH                                  | WHERE                                                         |
| WHERE                                  | FILTER                                                        |
| :label                                 | owl:Class or rdfs:Class                                       |
| [edge]                                 | Predicate <edge> or<br>Triple << <subject> <edge> <object> >> |
| (node)                                 | a graph node (blank or IRI node)                              |
| var:Person                             | ?var a :Person                                                |
| var.name                               | ?name assuming a match ?var <name> ?name                      |
| nodeVar:label {key value, key2 value2} | ?nodeVar a label; <key> value;key2 value2.                    |
| (ee)-[:KNOWS {since: 2001}]->(js)      | << ?ee <knows> <js> >> <since> 2001.                          |


RDF



# Grounding with Public Knowledge



GNOMON<sup>®</sup>  
DIGITAL



# Grounding with Public Knowledge




- Documents non structuré:
  - (textual similarity) TF-IDF / BM25
  - (dense encodings - semantic) Neural retriever AKA semantic similarity
- Knowledge Graphs
  - Graph traversal (e.g. n-hop)
  - Autoregressive path finding
  - GraphRag
- Web:
  - Off-the-shelf search engine for retrieval

# Grounding with Public Knowledge



Google search API , Elastic, MongoDB



# Grounding with Public Knowledge




GNOMON<sup>®</sup>  
DIGITAL

[https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex\\_genai/solutions/retrieval\\_augmented\\_generation.ipynb](https://github.com/GoogleCloudPlatform/asl-ml-immersion/blob/master/notebooks/vertex_genai/solutions/retrieval_augmented_generation.ipynb)

# Function Calling




La capacité de connecter de manière fiable les modèles de langage (LLMs) à des outils externes pour permettre une utilisation efficace de ces outils et une interaction avec des API externes.



# Function Calling



La capacité de connecter de manière fiable les modèles de langage (LLMs) à des outils externes pour permettre une utilisation efficace de ces outils et une interaction avec des API externes.







# Function Calling



La capacité de connecter de manière fiable les modèles de langage (LLMs) à des outils externes pour permettre une utilisation efficace de ces outils et une interaction avec des API externes.




# Function Calling



# Function Calling




GNOMON<sup>®</sup>  
DIGITAL



# Function Calling



GNOMON<sup>®</sup>  
DIGITAL



# Function Calling



GNOMON<sup>®</sup>  
DIGITAL

What is the weather like in London?

```
tools = [  
  {  
    "type": "function",  
    "function": {  
      "name": "get_current_weather",  
      "description": "Get the current weather in a given location",  
      "parameters": {  
        "type": "object",  
        "properties": {  
          "location": {  
            "type": "string",  
            "description": "The city and state, e.g. San Francisco, CA",  
          },  
          "unit": {  
            "type": "string",  
            "enum": ["celsius", "fahrenheit"]},  
        },  
        "required": ["location"],  
      },  
    },  
  },  
]
```

# Function Calling



GNOMON<sup>®</sup>  
DIGITAL

What is the weather like in London?

```
def get_completion(messages, model="gpt-3.5-turbo-1106", temperature=0, max_tokens=300, tools=None):  
    response = openai.chat.completions.create(  
        model=model,  
        messages=messages,  
        temperature=temperature,  
        max_tokens=max_tokens,  
        tools=tools  
    )  
    return response.choices[0].message
```



# Function Calling



GNOMON<sup>®</sup>  
DIGITAL

